Introduction

In statistics, the expectation–maximization (EM) algorithm is an iterative method for finding maximum likelihood or maximum a posteriori (MAP) estimates of parameters in statistical models, where the model depends on unobserved latent variables. For this reason EM is frequently used for data clustering, verification and identification of the speaker (biometric tasks), author profiling based on his documents, automatic document categorization, and many more applications.

em4gmm[1] is a toolkit to work with Finite Gaussian Mixture Models. In fact, it is a very fast and parallel C implementation of the clustering Expectation Maximization (EM) algorithm for estimating Gaussian Mixture Models (GMMs), with some extra important improvements. This is a full list of features that this toolkit contains:

- Fast learning of Gaussian Mixture Models (GMMs) using multidimensional data.

- Fast merge similar components to simplify the learned Gaussian Mixture Model.

- Fast simple classification score for a group of data (the data can be compressed).

- Obtain a detailed classification for each sample of a group of data (auto-clustering).

- Use a World Model to normalize classification scores (for biometric tasks).

- Possibility of set the number of threads to use on the execution of any task.

- Extra tool used to generate random synthetic data files, for any purpose.

In order to compile this toolkit, you need first compile and install the zlib library[2] from their website, or using your preferred software distribution channels (aptitude, yum, macports, etc) in order to install it (with the dev packages). On some systems this library can be installed by default.

Then, on Mac OS X and Linux distributions you can simple use the ***make*** command on the system shell to compile it, and then ***sudo make install*** to install it on yout system (by default on /usr/bin). We recommend the use of the lastest version of the GCC compiler (because the code generated by LLVM is, for now, much slower).

---

1   https://github.com/juandavm/em4gmm/
2   http://www.zlib.net/

Basic Usage

This toolkit has two different utilities: ***gmmtrain*** and ***gmmclass***. You can use the ***gmmtrain*** utility in order to train a Gaussian Mixture Model from a feature data file, and the ***gmmclass*** to classify (obtain the score/loglikelihood) one feature data file.

The data files used by this software are very simple: They are plain text files of decimal numbers, with a header, and one line per sample vector. This is a short example:

11   4

1025   7706   6830   5571   4169   2858   1809   1094   688   500   417

1147   5755   6636   6234   4118   4593   2750   3649   774   1568   1104

932   5381   5567   5175   3613   3499   2429   2536   652   913   337

838   6401   5961   5277   4418   3468   2516   1644   921   391   74

On the header, the first number are the dimension and the second the number of samples. The sample's vectors can be integers or decimals (using "." as separator), and the dimensions must be space-separated. Also, you have an example of a data file on the ***dat*** directory of this project. If you want to save disk space you can compress data files using gzip format (.gz file). A simple way to do this compression is using the gzip Linux or Mac Os X command.

You can train a Gaussian Mixture Model using the ***gmmtrain*** utility on a feature train file. This will learn the Model from the provided data, and save it into a file. This are all the options that can be used with this utility:

Usage: gmmtrain <options>

    Required:

        -d file.txt|file.gz   file that contains all the samples vectors

        -m file.gmm   file used to save the trained mixture model

    Recommended:

        -n 2-524228   optional number of components of the mixture

        -r file.json   optional file to save the log (not the model)

Optional:

> *-u 0.0-1.0*   *optional merge threshold based on similarity*
>
> *-s 0.0-1.0*   *optional stop criterion based on likelihood*
>
> *-i 1-10000*   *optional maximum number of EM iterations*
>
> *-t 1-256*   *optional maximum number of threads used*
>
> *-h*   *optional argument that shows this message*

Using the special *-u* parameter you can activate the merge feature that will try to merge the similar components of the mixture, leaving the model as simple as possible, avoiding overtraining, and speeding up the classifier. The parameter *-r* allows to obtain a log (it's a log and can not be used as the model to classify) of the generated model using the jSON text plain format. Also, yo can obtain the score/loglikelihood of one feature test file using the ***gmmclass*** utility:

Usage: gmmclass <options>

> *Required:*
>
> > *-d file.txt\file.gz*   *file that contains the samples vectors*
> >
> > *-m file.gmm*   *file of the trained model used to classify*
>
> *Recommended:*
>
> > *-w file.gmm*   *optional world model used to smooth*
> >
> > *-r file.json*   *optional file to save the log (slower)*
>
> *Optional:*
>
> > *-t 1-128*   *optional maximum number of threads used*
> >
> > *-h*   *optional argument that shows this message*

The standard process is to train a model for each class, and then classify at the class with highest score. Also, you can obtain a detailed analysis of the classify process done (for automatic clustering purposes, for instance) using the *-r* option on the ***gmmclass*** utility, but this analysis will run slower than the standard classification.

Synthetic Data Generation

We have included in this software a simple tool called ***datagen.py*** that can generate pseudo-random synthetic data using the normal distribution. This utility is used in this way:

Usage: datagen.py <dimensions> <samples> <file.txt>

As we can observe, use this tool is as easy as specify the number of samples, dimensions, and the file in will be saved the synthetic data. It is important to note that this tool is not as fast as the rest of the toolkit, because it was added as a simple extra tool.

The synthetic data file generated is fully compatible with the rest of the toolkit and can be used directly by it. The number of different normal distributions used to generate the data is given by the formula:

$$k = floor\left(\sqrt{\frac{n}{2}}\right) + 1$$

Where ***n*** is the number of samples and ***k*** the number of different normal distributions used to generate the final random synthetic data. The mean of the normal distribution is generated randomly, uniformly distributed in the range from -10000 to 10000. Also, the variance of the distributions is generated uniformly distributed in the range from 10 to 100.

This is an example of the result provided by this tool, using 5 dimensions and 8 samples:

5   7

-6139.657	*-9176.096*	*-9760.485*	*-2021.679*	*4101.922*
-6146.041	*-9189.280*	*-9812.855*	*-1966.770*	*4098.639*
-6128.852	*-9177.074*	*-9757.878*	*-1986.591*	*4126.399*
-6129.898	*-9144.684*	*-9714.653*	*-2000.045*	*4128.415*
8553.226	*7206.814*	*4854.733*	*-4044.405*	*-6667.375*
8595.694	*7259.777*	*4789.316*	*-4028.329*	*-6931.350*
8480.404	*7335.704*	*4869.384*	*-4069.297*	*-6675.915*

Speaker Identification

In order to make a speaker identification system, we need to record and extract acoustic features with any external toolkit (for instance the HTK[3] or TLK[4] toolkit). We need to create one Gaussian Mixture Model for each speaker, so first of all we put all the acoustic features for each speaker together on one file (with the required format for this toolkit, explained before).

Then, we need to train all the Gaussian Mixture Models using the ***gmmtrain*** utility of this toolkit for each speaker like this:

gmmtrain -d speakerN.feas -m speakerN.gmm -s 0.01 -u 0.95 -n 256

The ***-d*** parameter specify path to the acoustic features extracted before, the ***-m*** is the place to save the model, the ***-s*** parameter is a stop criterion based on the improvement, the ***-u*** allows to the toolkit merge similar components (this may be slow and use high memory, but avoids to find the optimal number of mixture components), and the ***-n*** is the initial number of components of the mixture. Of course, you can change the command to try to make the best model possible to your system, but this command probably can be a good startup point.

When you have all the models prepared, you can easy use them to identify the speaker. First of all the speaker must read or tell some short text, that must be recorded and transform into an acoustic feature like before (using the required format for this toolkit). The last step is, of course, identify the speaker. In order to do it, you will need to use the ***gmmclass*** utility of the toolkit in order to get the scores for each trained model. You can use it with a command like that:

gmmclass -d unknown.feas -m speakerN.gmm

The score for this feature will be printed on the standard output. The real speaker of these new recorded acoustic features are the speaker with the highest score, when we use this command with the same acoustic features on all the speaker models. It is important to note that the scores can be positive or negative, but in all cases the highest score is, probably, the real speaker (for instance, -1.97 is greater than -3.24, so the real speaker is the first).

This are the basics of the speaker identification task using the em4gmm toolkit.

---

3   http://htk.eng.cam.ac.uk/
4   http://www.translectures.eu/tlk/

Final Notes

Issues and Bugs

Do you have a bug or a feature request? Do not worry, open a new issue[5]. But please, before opening any new issue, search on existing the yours in order to avoid duplicates. And thanks you for your contribution!

Authors

Juan Daniel Valor Miró[6].

License

Expectation Maximization for Gaussian Mixture Models.

Copyright (C) 2012-2014 Juan Daniel Valor Miro.

This program is free software; you can redistribute it and/or

modify it under the terms of the GNU General Public License as

published by the Free Software Foundation; either version 2 of

the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful,

but WITHOUT ANY WARRANTY; without even the implied warranty of

MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

See the GNU General Public License for more details.

---

5   https://github.com/juandavm/em4gmm/issues/
6   http://www.juandaniel.es/